

1. Evaluate value of each State(i.e. $V(s)$) given below.

Value func $V(s) = E[G_t | S_t = s]$, return $G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$
 $= R_{t+1} + \gamma V(S_{t+1})$

$$V(s) = R_s + \gamma \sum_{s' \in S} P_{ss'} V(s')$$

$$\begin{bmatrix} V(\text{Sleep}) \\ V(\text{Pass}) \\ V(\text{fb}) \\ V(\text{Pub}) \\ V(\text{C1}) \\ V(\text{C2}) \\ V(\text{C3}) \end{bmatrix} = \begin{bmatrix} 0 \\ 10 \\ -1 \\ 1 \\ -2 \\ -2 \\ -2 \end{bmatrix} + \gamma \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0.9 & 0 & 0.1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0.2 & 0.4 & 0.4 \\ 0 & 0 & 0.5 & 0 & 0 & 0.5 & 0 \\ 0.2 & 0 & 0 & 0 & 0 & 0 & 0.8 \\ 0 & 0.6 & 0 & 0.4 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} V(\text{Sleep}) \\ V(\text{Pass}) \\ V(\text{fb}) \\ V(\text{Pub}) \\ V(\text{C1}) \\ V(\text{C2}) \\ V(\text{C3}) \end{bmatrix}$$

\parallel \parallel \parallel \parallel
 \mathbb{V} \mathbb{R} \mathbb{P} \mathbb{V}

GAMMA: 0	GAMMA: 0.9	GAMMA: 1
sleep: 0.00	sleep: 0.00	sleep: 0.00
pass: 10.00	pass: 10.00	pass: 10.00
facebook: -1.00	facebook: -7.64	facebook: -22.54
pub: 1.00	pub: 1.91	pub: 0.80
class 1: -2.00	class 1: -5.01	class 1: -12.54
class 2: -2.00	class 2: 0.94	class 2: 1.46
class 3: -2.00	class 3: 4.09	class 3: 4.32

\parallel \parallel \parallel
 $\gamma = 0$ $\gamma = 0.9$ $\gamma = 1$

```

1 import numpy as np
2
3 GAMMA = 1
4 states = {0:'sleep',
5           1:'pass',
6           2:'facebook',
7           3:'pub',
8           4:'class 1',
9           5:'class 2',
10          6:'class 3'}
11
12 with open('./rewards.txt', 'r') as f: # R
13     reward = f.readlines()
14
15 reward = np.array(reward, dtype=np.float64)
16
17 with open('./policy.txt', 'r') as f: # P
18     lines = f.readlines()
19
20 probs = []
21
22 for line in lines:
23     tmp = line.strip().split(' ')
24     probs.append(np.array(tmp, dtype=np.float64))
25
26 probs = np.array(probs)
27 I = np.eye(7)
28
29 inv_mat = np.linalg.inv(I-GAMMA*probs) # (I-gamma*P)
30 values = np.matmul(inv_mat,reward)     # (I-gamma*P)^-1 x R
31
32 for idx, val in enumerate(values):
33     print(f'{states[idx]:>8}: {val:<3.2f}')
34

```

```

rewards.txt
0
10
-1
1
-2
-2
-2

```

```


policy.txt
0 0 0 0 0 0 0
1 0 0 0 0 0 0
0 0 0.9 0 0.1 0 0
0 0 0 0 0.2 0.4 0.4
0 0 0.5 0 0 0.5 0
0.2 0 0 0 0 0 0.8
0 0.6 0 0.4 0 0 0

```

2. Obtain action values, $q(s,a)$ for each arrow in the given example.

$$q_{\pi}(s,a) = E_{\pi}[G_t | S_t=s, A_t=a]$$

$$q_{\pi}(s,a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \sum_{a' \in A} \pi(a'|s') q_{\pi}(s',a')$$



$$\begin{aligned}
q_{\pi}(\text{sleep}, \cdot) &= 0 & = 0 \\
q_{\pi}(\text{Pass}, \text{go2sleep}) &= 10 + \gamma \cdot 1.0 \cdot 0 & = 10 \\
q_{\pi}(\text{Class3}, \text{study}) &= -2 + \gamma \cdot 0.6 \cdot 10 & = 4 \\
q_{\pi}(\text{Class3}, \text{pub}) &= -2 + \gamma \cdot 0.4 \cdot V(\text{pub}) = -1.68 \\
q_{\pi}(\text{Class2}, \text{study}) &= -2 + \gamma \cdot 0.8 \cdot V(\text{Class3}) = 1.456 \\
q_{\pi}(\text{Class2}, \text{break}) &= -2 + \gamma \cdot 0.2 \cdot V(\text{sleep}) = -2 \\
q_{\pi}(\text{Class1}, \text{study}) &= -2 + \gamma \cdot 0.5 \cdot V(\text{Class2}) = -1.27 \\
q_{\pi}(\text{Class1}, \text{fb}) &= -2 + \gamma \cdot 0.5 \cdot V(\text{fb}) = -13.27 \\
q_{\pi}(\text{fb}, \text{fb}) &= -1 + \gamma \cdot 0.9 \cdot V(\text{fb}) = -21.286 \\
q_{\pi}(\text{fb}, \text{study}) &= -1 + \gamma \cdot 0.1 \cdot V(\text{Class1}) = -2.254 \\
q_{\pi}(\text{pub}, \text{studyP1}) &= 1 + \gamma \cdot 0.2 \cdot V(\text{Class1}) = -1.508 \\
q_{\pi}(\text{pub}, \text{studyP2}) &= 1 + \gamma \cdot 0.4 \cdot V(\text{Class2}) = 1.584 \\
q_{\pi}(\text{pub}, \text{studyP3}) &= 1 + \gamma \cdot 0.4 \cdot V(\text{Class3}) = 2.728
\end{aligned}$$

$$\gamma = 1$$

```

# action values
action_value = np.zeros(probs.shape)
for idx, state in enumerate(states.keys()):
    for i in range(len(list(states.keys()))):
        if probs[idx][i] != 0:
            action_value[idx][i] = reward[idx] + GAMMA*probs[idx][i]*values[i]
for idx, vec in enumerate(action_value):
    print(f'{states[idx]:>10}: ', end='')
    for element in vec:
        print(f'{element:^6.2f} ', end='')
    print()

```

GAMMA: 1

sleep: 0.00
 pass: 10.00
 facebook: -22.54
 pub: 0.80
 class 1: -12.54
 class 2: 1.46
 class 3: 4.32

sleep:	0.00	0.00	0.00	0.00	0.00	0.00	0.00
pass:	10.00	0.00	0.00	0.00	0.00	0.00	0.00
facebook:	0.00	0.00	-21.29	0.00	-2.25	0.00	0.00
pub:	0.00	0.00	0.00	0.00	-1.51	1.58	2.73
class 1:	0.00	0.00	-13.27	0.00	0.00	-1.27	0.00
class 2:	-2.00	0.00	0.00	0.00	0.00	0.00	1.46
class 3:	0.00	4.00	0.00	-1.68	0.00	0.00	0.00

break

$\gamma = 1$

pub

studyP1

studyP3

GAMMA: 0.9

sleep: 0.00
 pass: 10.00
 facebook: -7.64
 pub: 1.91
 class 1: -5.01
 class 2: 0.94
 class 3: 4.09

sleep:	0.00	0.00	0.00	0.00	0.00	0.00	0.00
pass:	10.00	0.00	0.00	0.00	0.00	0.00	0.00
facebook:	0.00	0.00	-7.19	0.00	-1.45	0.00	0.00
pub:	0.00	0.00	0.00	0.00	0.10	1.34	2.47
class 1:	0.00	0.00	-5.44	0.00	0.00	-1.58	0.00
class 2:	-2.00	0.00	0.00	0.00	0.00	0.00	0.94
class 3:	0.00	3.40	0.00	-1.31	0.00	0.00	0.00

$\gamma = 0.9$

3. How many iterations do we require to obtain the final value both for $V(s)$ and $q(s,a)$.

```
state_values = np.zeros(values.shape)
i = 0
while True:
    print(f'Iteration {i:>02}: ', end='')
    print_val(state_values)
    new_state_values = reward + np.matmul((GAMMA*probs), state_values)
    i += 1
    if np.power(np.power(state_values-new_state_values, 2), 0.5).sum() < 1e-3:
        state_values = new_state_values
        break
    state_values = new_state_values

print(f'Iteration {i:>02}: ', end='')
print_val(state_values)
print()
```

→ $V(S)$ 를 모두 0으로 두고
iterative하게 $V(S)$ 계산,

Iteration 00:	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Iteration 01:	0.00	10.00	-1.00	1.00	-2.00	-2.00	-2.00
Iteration 02:	0.00	10.00	-2.10	-1.00	-3.50	-3.60	4.40
Iteration 03:	0.00	10.00	-3.24	0.62	-4.85	1.52	3.60
Iteration 04:	0.00	10.00	-4.40	2.08	-2.86	0.88	4.25
Iteration 05:	0.00	10.00	-5.25	2.48	-3.76	1.40	4.83
Iteration 06:	0.00	10.00	-6.10	2.74	-3.92	1.86	4.99
Iteration 07:	0.00	10.00	-6.88	2.96	-4.12	1.99	5.10
Iteration 08:	0.00	10.00	-7.60	3.01	-4.44	2.08	5.18
Iteration 09:	0.00	10.00	-8.29	3.02	-4.76	2.15	5.20
Iteration 10:	0.00	10.00	-8.94	2.99	-5.07	2.16	5.21
Iteration 11:	0.00	10.00	-9.55	2.93	-5.39	2.16	5.20
Iteration 12:	0.00	10.00	-10.13	2.87	-5.69	2.16	5.17
Iteration 13:	0.00	10.00	-10.69	2.79	-5.99	2.14	5.15
Iteration 14:	0.00	10.00	-11.22	2.72	-6.28	2.12	5.12
Iteration 15:	0.00	10.00	-11.72	2.64	-6.55	2.09	5.09
Iteration 16:	0.00	10.00	-12.21	2.56	-6.82	2.07	5.06
Iteration 17:	0.00	10.00	-12.67	2.49	-7.07	2.04	5.02
Iteration 18:	0.00	10.00	-13.11	2.41	-7.31	2.02	4.99
Iteration 19:	0.00	10.00	-13.53	2.34	-7.54	2.00	4.97
Iteration 20:	0.00	10.00	-13.93	2.28	-7.77	1.97	4.94

⋮

Iteration 146:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 147:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 148:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 149:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 150:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 151:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 152:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 153:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 154:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 155:	0.00	10.00	-22.52	0.81	-12.53	1.46	4.32
Iteration 156:	0.00	10.00	-22.53	0.81	-12.53	1.46	4.32
Iteration 157:	0.00	10.00	-22.53	0.81	-12.53	1.46	4.32
Iteration 158:	0.00	10.00	-22.53	0.81	-12.53	1.46	4.32
Iteration 159:	0.00	10.00	-22.53	0.81	-12.53	1.46	4.32
Iteration 160:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32
Iteration 161:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32
Iteration 162:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32
Iteration 163:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32
Iteration 164:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32
Iteration 165:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32
Iteration 166:	0.00	10.00	-22.53	0.80	-12.54	1.46	4.32

→ 166 번째 iteration에서
작전 iter과 비교하여
 $1e-3$ 이하의 차이를 보임

Write algorithm code for a synchronous Value iteration agent.
Value iteration computes k-step estimates of the state values, $V(k)$ for $k=0, 1, 5, 10$ and ∞ .

1	2	3	+1 Goal with (award) 4
5	6	7	-1 Goal with (penalty) 8
9	10	11	12
13	14	15	16



- Initial policy : random
- Tabular MDP with 16 states
- Action : agent allow to move 4 normal direction
- Discount = none (gamma=1.0)
- reward = - 0.1 on all transition
- Two terminal states : one +1 award, the other -1 penalty


총 16개의 state 중 6, 7, 10, 11 은 상하좌우 모두 이동이 가능함

1, 13, 16은 2개의 방향으로만 이동가능

2, 3, 5, 9, 12, 14, 15는 8개의 방향으로 이동가능

4, 8은 terminal state로 이동 x

이 2 경우에 대해 움직일 수 없는 방향으로의 이동은 제자리로 돌아오도록함



< Code >

```

1 import numpy as np
2
3 def print_val(val):
4     for v in val:
5         print(f'{v:>5.2f}', end=' ')
6     print()
7
8 #코너에 있는 state들은 self-edge가 있는 case로 볼 수 있음
9 with open('./probs.csv', 'r', encoding='utf-8') as f:
10     lines = f.readlines()
11
12 probs = np.ndarray((16, 16))
13
14 for idx, line in enumerate(lines):
15     for jdx, p in enumerate(line.strip().split(',')):
16         probs[idx][jdx] = float(p)
17
18 reward = -0.1*np.ones(16)
19 reward[3] += 1
20 reward[7] += -1
21
22 GAMMA = 1.0
23
24 I = np.eye(16)
25
26 inv_mat = np.linalg.inv(I-GAMMA*probs)
27 values = np.matmul(inv_mat, reward)
28 for idx, val in enumerate(values):
29     print(f'{val:<3.2f}')
30 print(values.reshape(4, 4))
31
32 i = 0
33 state_values = np.zeros(reward.shape)
34 state_values[3] += 1
35 state_values[7] += -1
36 while True:
37     if i==0 or i == 1 or i == 5 or i == 10:
38         print(f'Iteration {i:>02}: ', end='')
39         print_val(state_values)
40         new_state_values = reward + np.matmul((GAMMA*probs), state_values)
41         i += 1
42         if np.abs(state_values-new_state_values).sum() < 1e-3:
43             state_values = new_state_values
44             break
45         state_values = new_state_values
46
47
48 print(f'Iteration {i:>02}: ', end='')
49 print_val(state_values)
50 print()
51
52 print(state_values.reshape(4, 4))

```

이동 -0.1 + 1 (award)
이동 -0.1 - 1 (penalty)

direct solution
iterative solution

Reward 값이 모두 0.1로 동일하므로 action에 대한 value의 max 값이 차이 나지 않음. (Terminal state에서는 다르게 설정했으나 전반적인 수렴의 경향에 큰 영향을 주지 않을 것으로 보여 생략함.)




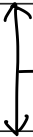










V k=0

0	0	0	+1
0	0	0	-1
0	0	0	0
0	0	0	0

			+1 Goal with (award)
			-1 Goal with (penalty)















V k=1

-0.1	-0.1	0.15	+1 0.9
-0.1	-0.1	-0.35	-1 -1.1
-0.1	-0.1	-0.1	-0.95
-0.1	-0.1	-0.1	-0.1

			+1 Goal with (award)
			-1 Goal with (penalty)
			
			















V k=5

-0.49	-0.4	-0.11	+1 0.9
-0.51	-0.53	-0.64	-1 -1.1
-0.52	-0.56	-0.65	-0.84
-0.51	-0.53	-0.6	-0.68

			+1 Goal with (award)
			-1 Goal with (penalty)
			
			










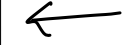




V k=10

-0.92	-0.78	-0.35	+1 Goal with 0.9 (award)
-0.98	-0.96	-0.93	-1 Goal with -1.1 (penalty)
-1.03	-1.05	-1.1	-1.19
-1.05	-1.08	-1.14	-1.20

			+1 Goal with (award)
			-1 Goal with (penalty)
			
			

V k=∞

-3.66	-3.05	-1.69	+1 Goal with 0.9 (award)
-3.89	-3.41	-2.49	-1 Goal with -1.1 (penalty)
-4.15	-3.84	-3.31	-2.95
-4.33	-4.11	-3.96	-3.45

			+1 Goal with (award)
			-1 Goal with (penalty)
			
			


```
Iteration 00:
[[ 0.  0.  0.  1.]
 [ 0.  0.  0. -1.]
 [ 0.  0.  0.  0.]
 [ 0.  0.  0.  0.]]
```

계산결과,,

```
Iteration 01:
[[-0.1 -0.1  0.15  0.9 ]
 [-0.1 -0.1 -0.35 -1.1 ]
 [-0.1 -0.1 -0.1  -0.35]
 [-0.1 -0.1 -0.1  -0.1 ]]
```

```
Iteration 05:
[[-0.47011719 -0.40117187 -0.10859375  0.9      ]
 [-0.5140625  -0.53027344 -0.64394531 -1.1      ]
 [-0.515625    -0.56367187 -0.65449219 -0.84082031]
 [-0.50976562 -0.53320312 -0.60273438 -0.68105469]]
```

```
Iteration 10:
[[-0.9153616  -0.77954102 -0.34628258  0.9      ]
 [-0.98342991 -0.95948486 -0.92506542 -1.1      ]
 [-1.03296738 -1.05198612 -1.1014637  -1.18898697]
 [-1.04966202 -1.08099174 -1.13628254 -1.19582729]]
```

```
Iteration 224:
[[-3.65751058 -3.04538685 -1.67207178  0.9      ]
 [-3.86995178 -3.4068396  -2.47098078 -1.1      ]
 [-4.1458376  -3.84131965 -3.30518746 -2.75307969]
 [-4.32659754 -4.1077307  -3.75561688 -3.45422419]]
```

probs

0.5	0.25	0	0	0.25	0	0	0	0	0	0	0	0	0	0	0
0.25	0.25	0.25	0	0	0.25	0	0	0	0	0	0	0	0	0	0
0	0.25	0.25	0.25	0	0	0.25	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0.25	0	0	0	0.25	0.25	0	0	0.25	0	0	0	0	0	0	0
0	0.25	0	0	0.25	0	0.25	0	0	0.25	0	0	0	0	0	0
0	0	0.25	0	0	0.25	0	0.25	0	0	0.25	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0.25	0	0	0	0.25	0.25	0	0	0.25	0	0	0
0	0	0	0	0	0.25	0	0	0.25	0	0.25	0	0	0.25	0	0
0	0	0	0	0	0	0.25	0	0	0.25	0	0.25	0	0	0.25	0
0	0	0	0	0	0	0	0.25	0	0	0.25	0.25	0	0	0	0.25
0	0	0	0	0	0	0	0	0	0.25	0	0	0	0.5	0.25	0
0	0	0	0	0	0	0	0	0	0	0.25	0	0	0.25	0.25	0.25
0	0	0	0	0	0	0	0	0	0	0	0.25	0	0	0.25	0.25
0	0	0	0	0	0	0	0	0	0	0	0	0.25	0	0	0.25

transition
probability 정의